

# NoHype: Virtualized Cloud Infrastructure without the Virtualization

Mehdi Nasef

Universidad de Cantabria

November 2023

# Introduction

# Introduction

- ▶ Virtualization only in software (With no help from hardware)

# Introduction

- ▶ Virtualization only in software (With no help from hardware)
- ▶ Virtualization in software with help from hardware

# Introduction

- ▶ Virtualization only in software (With no help from hardware)
- ▶ Virtualization in software with help from hardware
- ▶ NoHype: Virtualization only in hardware

# Motivation: Security concerns

## Motivation: Security concerns

- ▶ A successful attack on the hypervisor gives access to all VMs

# Motivation: Security concerns

- ▶ A successful attack on the hypervisor gives access to all VMs
- ▶ Large attack surface: Too many VM-Hypervisor interactions



# Motivation: Security concerns

- ▶ A successful attack on the hypervisor gives access to all VMs
- ▶ Large attack surface: Too many VM-Hypervisor interactions
- ▶ Securing the Hypervisor is becoming more and more difficult

# Threat model

# Threat model

- ▶ Cloud provider is not malicious



# Threat model

- ▶ Cloud provider is not malicious
- ▶ Guest operating systems are not trusted

# Threat model

- ▶ Cloud provider is not malicious
- ▶ Guest operating systems are not trusted
- ▶ Cloud management software is assumed to be secure
  - ▶ However, it is out of the scope of the paper.

# NoHype: Doing hypervisor functionality in hardware



# NoHype: Doing hypervisor functionality in hardware

## Controlling access to CPU

- ▶ Currently, VMs are scheduled like processes in a regular OS

# NoHype: Doing hypervisor functionality in hardware

## Controlling access to CPU

- ▶ Currently, VMs are scheduled like processes in a regular OS
- ▶ NoHype dedicates a core to only one VM



# NoHype: Doing hypervisor functionality in hardware

## Controlling access to CPU

- ▶ Currently, VMs are scheduled like processes in a regular OS
- ▶ NoHype dedicates a core to only one VM
  - ▶ It runs uninterrupted on it

# NoHype: Doing hypervisor functionality in hardware

## Controlling access to CPU

- ▶ Currently, VMs are scheduled like processes in a regular OS
- ▶ NoHype dedicates a core to only one VM
  - ▶ It runs uninterrupted on it
  - ▶ Multi-tenancy is still possible in many-core CPUs

# NoHype: Doing hypervisor functionality in hardware

## Controlling access to CPU

- ▶ Currently, VMs are scheduled like processes in a regular OS
- ▶ NoHype dedicates a core to only one VM
  - ▶ It runs uninterrupted on it
  - ▶ Multi-tenancy is still possible in many-core CPUs
  - ▶ VMs are more isolated this way

# NoHype: Doing hypervisor functionality in hardware

## Memory management

Hardware enforced memory partitioning:

# NoHype: Doing hypervisor functionality in hardware

## Memory management

Hardware enforced memory partitioning:

- ▶ Each VM has a dedicated share of physical memory

# NoHype: Doing hypervisor functionality in hardware

## Memory management

Hardware enforced memory partitioning:

- ▶ Each VM has a dedicated share of physical memory
- ▶ It gets mapped to its guest physical space when started

# NoHype: Doing hypervisor functionality in hardware

## Memory management

Hardware enforced memory partitioning:

- ▶ Each VM has a dedicated share of physical memory
- ▶ It gets mapped to its guest physical space when started
- ▶ The guest OS manages it as it would do without virtualization

# NoHype: Doing hypervisor functionality in hardware

## Memory management

Hardware enforced memory partitioning:

- ▶ Each VM has a dedicated share of physical memory
- ▶ It gets mapped to its guest physical space when started
- ▶ The guest OS manages it as it would do without virtualization
- ▶ Hardware support is already present on current machines:  
Extended Page Tables





# NoHype: Doing hypervisor functionality in hardware

## Memory management

Memory bandwidth fairness:

# NoHype: Doing hypervisor functionality in hardware

## Memory management

Memory bandwidth fairness:

- ▶ Prevent VMs from hogging memory bandwidth

# NoHype: Doing hypervisor functionality in hardware

## Memory management

Memory bandwidth fairness:

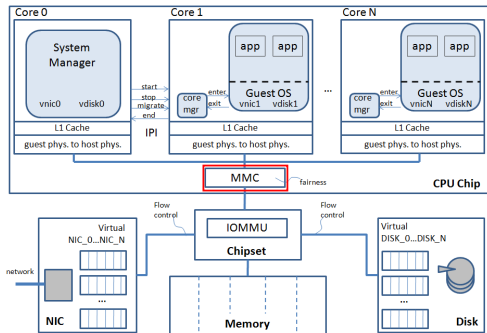
- ▶ Prevent VMs from hogging memory bandwidth
- ▶ Use a fair memory request scheduling algorithm

# NoHype: Doing hypervisor functionality in hardware

## Memory management

Memory bandwidth fairness:

- ▶ Prevent VMs from hogging memory bandwidth
- ▶ Use a fair memory request scheduling algorithm
- ▶ It is not supported by current hardware



# NoHype: Doing hypervisor functionality in hardware

IO devices

# NoHype: Doing hypervisor functionality in hardware

## IO devices

- ▶ Dedicated actual or hardware virtual IO devices per VM (direct access)

# NoHype: Doing hypervisor functionality in hardware

## IO devices

- ▶ Dedicated actual or hardware virtual IO devices per VM (direct access)
- ▶ Device interfaces are mapped to the VM's guest physical space



# NoHype: Doing hypervisor functionality in hardware

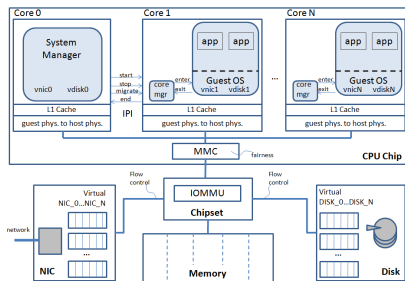
## IO devices

- ▶ Dedicated actual or hardware virtual IO devices per VM (direct access)
- ▶ Device interfaces are mapped to the VM's guest physical space
- ▶ control the rate of IO operations: PCIe flow control



# Start, stop and abort a VM

Some software is still needed for these functions



## Start a VM

1. Cloud manager notifies system manager to start a VM
2. System manager maps the memory and disk of the to-be-assigned VM into the system manager's space
3. System manager downloads the disk image and zeroes out the memory
4. After initialization, system manager un-maps the memory and disk from its space
5. System manager issues a 'start' IPI to the core
6. Core manager initializes memory and I/O mapping
7. Core manager exits to guest OS, starting the guest OS execution

## Stop a VM

1. Cloud manager notifies system manager to stop a VM
2. System manager issues a 'stop' IPI to the core running the target VM
3. Core manager optionally saves the disk image of the VM, then clears the disk and memory
4. Core manager un-maps memory and I/O
5. Core manager puts the current core into idle state

## Abort a VM

1. Enter core manager on an illegal operation
2. Core manager sends an 'end' IPI to the system manager
3. Core manager optionally clears disk and memory
4. Core manager un-maps memory and I/O
5. Core manager puts the current core into idle state
6. System manager notifies cloud manager

# Conclusions

Thank you for your attention!